

Spectral Neural Networks: A Universal Approximation Framework via Spectral Decomposition

Orges Leka

November 1, 2024

Abstract

Neural network architectures that draw inspiration from physical interactions have garnered significant interest due to their structured parameterizations and potential computational efficiencies. This paper introduces the concept of *Spectral Neural Networks*, where the weight matrices are parameterized through spectral (eigen-decomposition) representations. We investigate the representational capacity of such networks, particularly focusing on their ability to universally approximate any symmetric weight matrix. By leveraging the spectral theorem, we establish conditions under which spectral neural networks achieve universal approximation. Furthermore, we discuss the implications of dimensionality constraints, propose strategies to enhance the expressiveness of these networks under limited spectral dimensions, and elucidate the learning mechanisms involving gradient descent and backpropagation within this framework.

1 Introduction

The universal approximation capability of neural networks underpins their widespread applicability in diverse machine learning tasks. Traditional architectures, such as multilayer perceptrons (MLPs), achieve this by employing dense weight matrices that can represent arbitrary linear transformations given sufficient capacity. However, as the scale of neural networks grows, the parameterization of weight matrices becomes a computational bottleneck, motivating the exploration of more structured approaches.

In this context, we explore a neural network architecture inspired by physical interactions, where each neuron is associated with "mass" and "position" vectors. The weight matrix between neurons is computed using a force-like inverse-square law interaction, analogous to the Coulomb interaction between charged particles. Specifically, for neurons with mass vectors μ_i and μ_j located at positions x_i and x_j , the weight w_{ij} is defined as:

$$w_{ij} = \frac{\mu_i \cdot \mu_j}{\|x_i - x_j\|^2}$$

This formulation bears structural resemblance to the *Coulomb matrix* used in quantum chemistry to represent atomic interactions in molecules.

The central inquiry of this paper is to determine the conditions under which a general symmetric weight matrix W can be represented in the form of a Coulomb-like matrix. Formally, we seek to establish when there exists a set of vectors $\{\mu_i, x_i\}$ such that:

$$w_{ij} = \frac{\mu_i \cdot \mu_j}{\|x_i - x_j\|^2} \quad \forall i, j$$

Subsequently, we refine the problem to the case where the diagonal of W is zero:

$$w_{ij} = \mu_i \cdot \mu_j \cdot \|x_i - x_j\|^{-2} \quad \forall i \neq j$$

This simplification aids in the analytical tractability of the representation.

2 Spectral Neural Networks

2.1 Spectral Parameterization of Weight Matrices

A *Spectral Neural Network* parameterizes its weight matrices via spectral decomposition. Specifically, a symmetric weight matrix $W \in \mathbb{R}^{n \times n}$ is expressed as:

$$W = Q\Lambda Q^\top$$

where:

- $Q \in \mathbb{R}^{n \times d}$ is a matrix whose columns are vectors q_k , analogous to eigenvectors.
- $\Lambda \in \mathbb{R}^d$ is a vector containing scaling factors λ_k , analogous to eigenvalues.
- d is the dimensionality parameter, potentially much smaller than n , the number of neurons.

This formulation draws inspiration from the *Spectral Theorem*, which asserts that any real symmetric matrix can be decomposed into its eigenvectors and eigenvalues.

2.2 Relation to Coulomb-like Matrices

In the original formulation, weights are defined via interactions resembling Coulomb forces:

$$w_{ij} = \frac{\mu_i \cdot \mu_j}{\|x_i - x_j\|^2}$$

However, for analytical simplicity, we consider the modified representation without division by the distance squared:

$$w_{ij} = \mu_i \cdot \mu_j \cdot \|x_i - x_j\|^2 \quad \forall i \neq j$$

This adjustment facilitates the exploration of the representational capacity without singularities arising from zero distances.

3 Universal Approximation with Spectral Neural Networks

3.1 Universal Approximation Theorem

Theorem 1. A *Spectral Neural Network* can universally approximate any symmetric weight matrix $W \in \mathbb{R}^{n \times n}$ if and only if the dimensionality parameter d satisfies $d \geq n$.

3.2 Proof of Theorem 1

Sufficiency ($d \geq n$):

By the *Spectral Theorem*, any real symmetric matrix $W \in \mathbb{R}^{n \times n}$ can be decomposed as:

$$W = Q\Lambda Q^\top$$

where Q is an orthogonal matrix ($Q^\top Q = I$) and Λ is a diagonal matrix containing the eigenvalues of W .

When $d \geq n$, $Q \in \mathbb{R}^{n \times n}$ can encompass all the eigenvectors of W , and $\Lambda \in \mathbb{R}^n$ can include all corresponding eigenvalues. Thus, the spectral parameterization can exactly represent any symmetric weight matrix W .

Necessity ($d < n$):

If $d < n$, the matrix $Q\Lambda Q^\top$ has rank at most d , since it is the product of an $n \times d$ matrix, a d -dimensional diagonal matrix, and a $d \times n$ matrix. Consequently, only symmetric matrices of rank $\leq d$ can be exactly represented. However, there exist symmetric matrices with rank $> d$, which cannot be represented by the spectral parameterization when $d < n$. Therefore, exact universal approximation is unattainable in this regime.

3.3 Implications of Theorem 1

The theorem establishes that spectral neural networks achieve universal approximation of symmetric weight matrices when the spectral dimensionality meets or exceeds the number of neurons. However, practical scenarios often require $d \ll n$ for computational efficiency, necessitating approximate representations through low-rank decompositions.

4 Learning in Spectral Neural Networks

4.1 Network Parameterization

In the Spectral Neural Network framework, the entire network's weight matrices are parameterized through spectral components. Specifically, each weight matrix W in the network is represented as:

$$W^{(l)} = Q^{(l)} \Lambda^{(l)} (Q^{(l)})^\top$$

for each layer l , where:

- $Q^{(l)} \in \mathbb{R}^{n_l \times d_l}$ contains the spectral vectors $q_k^{(l)}$ for layer l .
- $\Lambda^{(l)} \in \mathbb{R}^{d_l}$ is the scaling vector $\lambda_k^{(l)}$ for layer l .
- d_l is the spectral dimensionality for layer l .

Each neuron within the network is associated with a vector $q_i^{(l)}$ in $Q^{(l)}$, encapsulating its spectral properties. The scaling factors $\Lambda^{(l)}$ modulate the contributions of each spectral component to the overall weight matrix.

4.2 Learning Mechanism

The learning process in Spectral Neural Networks involves optimizing the spectral parameters $Q^{(l)}$ and $\Lambda^{(l)}$ across all layers. This optimization is performed using gradient descent in conjunction with backpropagation to minimize a predefined loss function L .

4.2.1 Gradient Descent

Given a loss function L , the objective is to minimize L with respect to the spectral parameters. The update rules for gradient descent are as follows:

$$Q^{(l)} \leftarrow Q^{(l)} - \eta \frac{\partial L}{\partial Q^{(l)}}$$

$$\Lambda^{(l)} \leftarrow \Lambda^{(l)} - \eta \frac{\partial L}{\partial \Lambda^{(l)}}$$

where η is the learning rate.

4.2.2 Backpropagation

Backpropagation is employed to compute the gradients $\frac{\partial L}{\partial Q^{(l)}}$ and $\frac{\partial L}{\partial \Lambda^{(l)}}$ efficiently. The computation leverages the chain rule to propagate errors from the output layer backward through each spectral parameterized layer.

For each layer l , the gradient with respect to $Q^{(l)}$ is derived as:

$$\frac{\partial L}{\partial Q^{(l)}} = \frac{\partial L}{\partial W^{(l)}} \cdot \frac{\partial W^{(l)}}{\partial Q^{(l)}} = \frac{\partial L}{\partial W^{(l)}} \cdot \left(\Lambda^{(l)} (Q^{(l)})^\top + Q^{(l)} (\Lambda^{(l)})^\top \right)$$

Similarly, the gradient with respect to $\Lambda^{(l)}$ is:

$$\frac{\partial L}{\partial \Lambda^{(l)}} = \text{diag} \left(Q^{(l)\top} \cdot \frac{\partial L}{\partial W^{(l)}} \cdot Q^{(l)} \right)$$

where $\text{diag}(\cdot)$ extracts the diagonal elements corresponding to the scaling factors.

4.3 Parameter Efficiency

By parameterizing weight matrices through spectral decomposition, Spectral Neural Networks can achieve significant parameter efficiency, especially when $d_l \ll n_l$ for each layer l . This reduction from $O(n_l^2)$ to $O(n_l d_l)$ parameters not only conserves memory but also potentially enhances learning dynamics by focusing on the most salient spectral components.

5 Enhancing Expressiveness with Limited Spectral Dimensionality

While universal approximation is guaranteed for $d \geq n$, practical applications benefit from lower dimensionality. To mitigate the expressiveness limitations when $d < n$, the following strategies are proposed:

5.1 Hierarchical Spectral Decomposition

Introduce multiple layers of spectral decompositions within the network architecture. Each layer captures distinct aspects or frequency components of the transformations, thereby compensating for the reduced dimensionality in individual layers.

5.2 Sparsity Constraints

Impose sparsity on the matrices $Q^{(l)}$ or the vectors $\Lambda^{(l)}$. Sparse representations focus on the most significant interactions, effectively utilizing limited dimensionality to capture essential patterns within the weight matrices.

5.3 Nonlinear Transformations

Incorporate nonlinear activation functions between spectral layers. Nonlinearities enable the network to model complex interactions and compensate for the expressiveness loss due to reduced spectral dimensionality.

5.4 Adaptive Dimensionality

Allow the spectral dimensionality parameter d_l to adapt dynamically during training. Starting with a small d_l , the network can incrementally increase dimensionality based on the complexity of the task, balancing computational efficiency with representational capacity.

6 Conclusion

This paper introduces Spectral Neural Networks, a novel architecture wherein symmetric weight matrices are parameterized via spectral decomposition. We established that such networks achieve universal approximation of symmetric weight matrices provided the spectral dimensionality meets or exceeds the number of neurons. While exact universality necessitates $d \geq n$, practical applications can leverage strategies such as hierarchical decomposition, sparsity, nonlinear

transformations, and adaptive dimensionality to enhance expressiveness under computational constraints.

Furthermore, we elucidated the learning mechanisms within Spectral Neural Networks, highlighting how gradient descent and backpropagation facilitate the optimization of spectral parameters. By structuring the network’s parameters through λ vectors and q vectors for each neuron, the framework not only ensures parameter efficiency but also integrates seamlessly with established optimization techniques in neural network training.

The spectral parameterization offers a structured and potentially memory-efficient representation of weight matrices, bridging concepts from physics-inspired interactions and linear algebraic decompositions. Future work may explore empirical validations of the proposed theoretical insights and the integration of spectral methods with other architectural innovations to further optimize performance and scalability.

References

1. Horn, R. A., & Johnson, C. R. (2012). *Matrix Analysis*. Cambridge University Press.
2. Hornik, K., Stinchcombe, M., & White, H. (1989). Multilayer feedforward networks are universal approximators. *Neural Networks*, 2(5), 359-366.
3. Saxe, A. M., McClelland, J. L., & Ganguli, S. (2013). Exact solutions to the nonlinear dynamics of learning in deep linear neural networks. *arXiv preprint arXiv:1312.6120*.
4. Candes, E. J., & Recht, B. (2009). Exact matrix completion via convex optimization. *Foundations of Computational Mathematics*, 9(6), 717-772.
5. Coulomb, C. A. (1785). *Recherches sur la théorie des phénomènes électro-dynamiques*.

A Appendix: Python Implementation of Spectral Neural Networks

Below is the Python code used for implementing and experimenting with Spectral Neural Networks. The code includes the definition of activation functions, loss functions, the spectral layer, the spectral neural network class, and example usage on XOR, classification, and regression tasks.

```

1 import numpy as np
2 from sklearn.datasets import load_breast_cancer, load_diabetes
3 from sklearn.model_selection import train_test_split
4 from sklearn.preprocessing import StandardScaler
5 from sklearn.metrics import (
6     accuracy_score,
7     mean_squared_error,
8     matthews_corrcoef,
9     r2_score,
10    confusion_matrix,
11    roc_curve,
12    auc
13 )
14 from sklearn.svm import SVC
15 from sklearn.neural_network import MLPClassifier
16 from sklearn.linear_model import LogisticRegression, LinearRegression
17 from sklearn.ensemble import RandomForestClassifier, RandomForestRegressor
18 import matplotlib.pyplot as plt
19 import seaborn as sns
20

```

```

21 # Activation functions and their derivatives
22 def relu(x):
23     return np.maximum(0, x)
24
25 def relu_derivative(x):
26     return (x > 0).astype(float)
27
28 def sigmoid(x):
29     return 1 / (1 + np.exp(-x))
30
31 def sigmoid_derivative(x):
32     s = sigmoid(x)
33     return s * (1 - s)
34
35 # Loss functions and their derivatives
36 def mse_loss(y_true, y_pred):
37     return np.mean((y_true - y_pred) ** 2)
38
39 def mse_loss_derivative(y_true, y_pred):
40     return 2 * (y_pred - y_true) / y_true.size
41
42 def binary_cross_entropy(y_true, y_pred):
43     # Adding epsilon for numerical stability
44     epsilon = 1e-12
45     y_pred = np.clip(y_pred, epsilon, 1 - epsilon)
46     return -np.mean(y_true * np.log(y_pred) + (1 - y_true) * np.log(1 - y_pred)
47 )
48
49 def binary_cross_entropy_derivative(y_true, y_pred):
50     # Adding epsilon for numerical stability
51     epsilon = 1e-12
52     y_pred = np.clip(y_pred, epsilon, 1 - epsilon)
53     return (-(y_true / y_pred) + (1 - y_true) / (1 - y_pred)) / y_true.size
54
55 # Spectral Layer Class for Rectangular Weight Matrices
56 class SpectralLayer:
57     """
58     Initializes the Spectral Layer.
59
60     Parameters:
61     - input_dim: Number of input neurons.
62     - output_dim: Number of output neurons.
63     - spectral_dim: Spectral dimensionality 'd'.
64     - activation: Activation function ('relu' or 'sigmoid').
65     """
66     self.input_dim = input_dim
67     self.output_dim = output_dim
68     self.spectral_dim = spectral_dim
69
70     # Initialize Q and P with He initialization for better convergence
71     # Q: (output_dim, d)
72     self.Q = np.random.randn(output_dim, spectral_dim) * np.sqrt(2. / (
73 output_dim + spectral_dim))
74     # P: (input_dim, d)
75     self.P = np.random.randn(input_dim, spectral_dim) * np.sqrt(2. / (
76 input_dim + spectral_dim))
77     # Lambda: (d,)
78     self.Lambda = np.random.randn(spectral_dim) * 0.1
79
80     # Initialize biases
81     self.b = np.zeros((output_dim,))

```

```

81     # Activation function
82     if activation == 'relu':
83         self.activation = relu
84         self.activation_derivative = relu_derivative
85     elif activation == 'sigmoid':
86         self.activation = sigmoid
87         self.activation_derivative = sigmoid_derivative
88     else:
89         raise ValueError("Unsupported activation function")
90
91     # Placeholders for forward and backward pass
92     self.x = None
93     self.z = None
94     self.a = None
95     self.dQ = None
96     self.dP = None
97     self.dLambda = None
98     self.db = None
99
100 def forward(self, x):
101     """
102     Forward pass through the spectral layer.
103
104     Parameters:
105     - x: Input data of shape (batch_size, input_dim)
106
107     Returns:
108     - a: Activated output of shape (batch_size, output_dim)
109     """
110     self.x = x # (batch_size, input_dim)
111
112     # Compute  $W = Q * \text{diag}(\text{Lambda}) * P.T$ 
113     # Q_Lambda: (output_dim, d)
114     Q_Lambda = self.Q * self.Lambda # Broadcasting (output_dim, d)
115     # Compute  $W = Q\_Lambda @ P.T \rightarrow (\text{output\_dim}, \text{input\_dim})$ 
116     self.W = Q_Lambda @ self.P.T # (output_dim, input_dim)
117
118     # Compute linear transformation
119     self.z = self.W @ x.T + self.b[:, np.newaxis] # (output_dim,
batch_size)
120
121     # Apply activation
122     self.a = self.activation(self.z) # (output_dim, batch_size)
123
124     return self.a.T # (batch_size, output_dim)
125
126 def backward(self, delta):
127     """
128     Backward pass through the spectral layer.
129
130     Parameters:
131     - delta: Gradient of loss with respect to activated output (batch_size,
output_dim)
132
133     Returns:
134     - delta_prev: Gradient of loss with respect to input x (batch_size,
input_dim)
135     """
136     batch_size = self.x.shape[0]
137
138     # Compute derivative of activation
139     dz = delta.T * self.activation_derivative(self.z) # (output_dim,
batch_size)

```

```

140
141     # Compute gradients w.r. to biases
142     self.db = np.sum(dz, axis=1) / batch_size # (output_dim,)
143
144     # Compute gradients w.r. to W
145     # W = Q * diag(Lambda) * P.T
146     # dL/dW = dz @ x / batch_size
147     dW = dz @ self.x / batch_size # (output_dim, input_dim)
148
149     # Initialize gradients
150     self.dQ = np.zeros_like(self.Q) # (output_dim, d)
151     self.dP = np.zeros_like(self.P) # (input_dim, d)
152     self.dLambda = np.zeros_like(self.Lambda) # (d,)
153
154     # Compute gradients w.r. to Q, P, and Lambda
155     for k in range(self.spectral_dim):
156         Q_k = self.Q[:, k].reshape(-1, 1) # (output_dim, 1)
157         P_k = self.P[:, k].reshape(-1, 1) # (input_dim, 1)
158         Lambda_k = self.Lambda[k]
159
160         # Gradient w.r. Q_k: (output_dim, 1) += Lambda_k * (dW @ P_k)
161         self.dQ[:, k:k+1] += Lambda_k * (dW @ P_k) # (output_dim, 1)
162
163         # Gradient w.r. P_k: (input_dim, 1) += Lambda_k * (dW.T @ Q_k)
164         self.dP[:, k:k+1] += Lambda_k * (dW.T @ Q_k) # (input_dim, 1)
165
166         # Gradient w.r. Lambda_k: sum of element-wise product
167         self.dLambda[k] += np.sum(dW * (Q_k @ P_k.T)) # Scalar
168
169     # Compute gradient w.r. to input x
170     # dL/dx = W.T @ dz
171     delta_prev = self.W.T @ dz # (input_dim, batch_size)
172     return delta_prev.T # (batch_size, input_dim)
173
174 def update_parameters(self, learning_rate):
175     """
176     Updates the spectral parameters Q, Lambda, P, and biases using gradient
177     descent.
178
179     Parameters:
180     - learning_rate: Learning rate for gradient descent.
181     """
182     self.Q -= learning_rate * self.dQ
183     self.P -= learning_rate * self.dP
184     self.Lambda -= learning_rate * self.dLambda
185     self.b -= learning_rate * self.db
186
187 # Spectral Neural Network Class
188 class SpectralNeuralNetwork:
189     def __init__(self, layer_dims, spectral_dims, activations):
190         """
191         Initializes the Spectral Neural Network.
192
193         Parameters:
194         - layer_dims: List of neuron counts for each layer, including input and
195         output layers.
196         - spectral_dims: List of spectral dimensions for each layer (excluding
197         input layer).
198         - activations: List of activation functions for each layer (excluding
199         input layer).
200         """
201         assert len(layer_dims) - 1 == len(spectral_dims) == len(activations), "
Mismatch in layer specifications."

```



```

198     self.layers = []
199     for i in range(len(layer_dims) - 1):
200         layer = SpectralLayer(
201             input_dim=layer_dims[i],
202             output_dim=layer_dims[i+1],
203             spectral_dim=spectral_dims[i],
204             activation=activations[i]
205         )
206         self.layers.append(layer)
207     self.loss_history = [] # To store loss values during training
208
209     def forward(self, x):
210         """
211         Forward pass through the network.
212
213         Parameters:
214         - x: Input data of shape (batch_size, input_dim)
215
216         Returns:
217         - Output of the network
218         """
219         for layer in self.layers:
220             x = layer.forward(x)
221         return x
222
223     def backward(self, loss_grad):
224         """
225         Backward pass through the network.
226
227         Parameters:
228         - loss_grad: Gradient of the loss with respect to the network's output
229         """
230         for layer in reversed(self.layers):
231             loss_grad = layer.backward(loss_grad)
232
233     def update_parameters(self, learning_rate):
234         """
235         Updates all spectral layers' parameters.
236
237         Parameters:
238         - learning_rate: Learning rate for gradient descent.
239         """
240         for layer in self.layers:
241             layer.update_parameters(learning_rate)
242
243     def train(self, X, y, epochs, learning_rate, loss_function='mse'):
244         """
245         Trains the network using gradient descent.
246
247         Parameters:
248         - X: Training data of shape (num_samples, input_dim)
249         - y: Training labels of shape (num_samples, output_dim)
250         - epochs: Number of training epochs
251         - learning_rate: Learning rate for gradient descent
252         - loss_function: 'mse' or 'bce'
253         """
254         for epoch in range(epochs):
255             # Forward pass
256             output = self.forward(X) # (num_samples, output_dim)
257
258             # Compute loss
259             if loss_function == 'mse':
260                 loss = mse_loss(y, output)

```

```

261         loss_grad = mse_loss_derivative(y, output) # (num_samples,
output_dim)
262         elif loss_function == 'bce':
263             loss = binary_cross_entropy(y, output)
264             loss_grad = binary_cross_entropy_derivative(y, output) # (
num_samples, output_dim)
265         else:
266             raise ValueError("Unsupported loss function")
267
268         self.loss_history.append(loss)
269
270         # Backward pass
271         self.backward(loss_grad)
272
273         # Update parameters
274         self.update_parameters(learning_rate)
275
276         # Print loss every 10% of epochs or first epoch
277         if (epoch + 1) % (epochs // 10) == 0 or epoch == 0:
278             print(f"Epoch {epoch+1}/{epochs}, Loss: {loss:.4f}")
279
280     def predict(self, X):
281         """
282         Makes predictions with the network.
283
284         Parameters:
285         - X: Input data of shape (num_samples, input_dim)
286
287         Returns:
288         - Predictions of shape (num_samples, output_dim)
289         """
290         return self.forward(X)
291
292 # Example Usage: XOR, Classification, and Regression Problems
293 def create_xor_data():
294     """
295     Creates XOR dataset.
296
297     Returns:
298     - X: Input data of shape (4, 2)
299     - y: Labels of shape (4, 1)
300     """
301     X = np.array([
302         [0, 0],
303         [0, 1],
304         [1, 0],
305         [1, 1]
306     ])
307     y = np.array([
308         [0],
309         [1],
310         [1],
311         [0]
312     ])
313     return X, y
314
315 def plot_loss(history, title):
316     """
317     Plots the loss curve.
318
319     Parameters:
320     - history: List of loss values.
321     - title: Title of the plot.

```

```

322     """
323     plt.figure(figsize=(8,6))
324     plt.plot(history, label='Loss')
325     plt.title(title)
326     plt.xlabel('Epochs')
327     plt.ylabel('Loss')
328     plt.legend()
329     plt.grid(True)
330     plt.show()
331
332 def plot_classification_comparison(mcc_scores, models, title):
333     """
334     Plots a bar chart comparing MCC scores across models.
335
336     Parameters:
337     - mcc_scores: List of MCC scores.
338     - models: List of model names.
339     - title: Title of the plot.
340     """
341     plt.figure(figsize=(10,6))
342     sns.barplot(x=models, y=mcc_scores, palette='viridis')
343     plt.title(title)
344     plt.ylabel('Matthews Correlation Coefficient (MCC)')
345     plt.ylim(-1,1)
346     plt.xticks(rotation=45)
347     plt.show()
348
349 def plot_regression_comparison(r2_scores, models, title):
350     """
351     Plots a bar chart comparing R2 scores across models.
352
353     Parameters:
354     - r2_scores: List of R2 scores.
355     - models: List of model names.
356     - title: Title of the plot.
357     """
358     plt.figure(figsize=(10,6))
359     sns.barplot(x=models, y=r2_scores, palette='magma')
360     plt.title(title)
361     plt.ylabel('R2 Score')
362     plt.ylim(0,1)
363     plt.xticks(rotation=45)
364     plt.show()
365
366 def plot_confusion_matrix(cm, classes, title):
367     """
368     Plots the confusion matrix.
369
370     Parameters:
371     - cm: Confusion matrix.
372     - classes: List of class names.
373     - title: Title of the plot.
374     """
375     plt.figure(figsize=(6,5))
376     sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
377                 xticklabels=classes, yticklabels=classes)
378     plt.ylabel('Actual')
379     plt.xlabel('Predicted')
380     plt.title(title)
381     plt.show()
382
383 def plot_roc_curve(y_true, y_scores, title):
384     """

```

```

385     Plots the ROC curve.
386
387     Parameters:
388     - y_true: True binary labels.
389     - y_scores: Scores/probabilities from the classifier.
390     - title: Title of the plot.
391     """
392     fpr, tpr, thresholds = roc_curve(y_true, y_scores)
393     roc_auc = auc(fpr, tpr)
394
395     plt.figure(figsize=(8,6))
396     plt.plot(fpr, tpr, color='darkorange',
397              lw=2, label=f'ROC curve (area = {roc_auc:.2f})')
398     plt.plot([0,1], [0,1], color='navy', lw=2, linestyle='--')
399     plt.xlim([-0.01,1.0])
400     plt.ylim([0.0,1.05])
401     plt.xlabel('False Positive Rate')
402     plt.ylabel('True Positive Rate')
403     plt.title(title)
404     plt.legend(loc="lower right")
405     plt.grid(True)
406     plt.show()
407
408 def plot_regression_predictions(y_true, y_pred, title):
409     """
410     Plots predicted vs actual values for regression.
411
412     Parameters:
413     - y_true: True target values.
414     - y_pred: Predicted target values.
415     - title: Title of the plot.
416     """
417     plt.figure(figsize=(8,6))
418     plt.scatter(y_true, y_pred, alpha=0.7, label='Spectral NN')
419     plt.plot([y_true.min(), y_true.max()], [y_true.min(), y_true.max()], 'r--',
420              lw=2, label='Ideal Fit')
421     plt.xlabel('Actual Values')
422     plt.ylabel('Predicted Values')
423     plt.title(title)
424     plt.legend()
425     plt.grid(True)
426     plt.show()
427
428 def main():
429     # =====
430     # Part 1: XOR Problem
431     # =====
432     print("Training on XOR Problem")
433     X_xor, y_xor = create_xor_data()
434
435     # Define network architecture for XOR
436     input_dim = 2
437     hidden_dim = 4
438     output_dim = 1
439     spectral_dim_hidden = 7 # Set spectral_dim >= hidden_dim for better
440                             expressiveness
441     spectral_dim_output = 1 # Corrected spectral_dim_output to match
442                             output_dim
443
444     layer_dims_xor = [input_dim, hidden_dim, output_dim]
445     spectral_dims_xor = [spectral_dim_hidden, spectral_dim_output]
446     activations_xor = ['relu', 'sigmoid']

```

```

445 # Initialize the network
446 network_xor = SpectralNeuralNetwork(layer_dims_xor, spectral_dims_xor,
activations_xor)
447
448 # Train the network on XOR
449 epochs = 10000
450 learning_rate = 0.1 # Increased learning rate for faster convergence
451 network_xor.train(X_xor, y_xor, epochs, learning_rate, loss_function='bce')
452
453 # Plot loss curve for XOR
454 plot_loss(network_xor.loss_history, "Spectral Neural Network Loss Curve for
XOR Problem")
455
456 # Make predictions on XOR
457 predictions_xor = network_xor.predict(X_xor)
458 predictions_binary_xor = (predictions_xor > 0.5).astype(int)
459
460 print("\nPredictions on XOR after training:")
461 for i in range(len(X_xor)):
462     print(f"Input: {X_xor[i]}, Predicted: {predictions_binary_xor[i][0]},
True: {y_xor[i][0]}")
463
464 # =====
465 # Part 2: Binary Classification on Breast Cancer Dataset
466 # =====
467 print("\n\nTraining on Breast Cancer Classification")
468
469 # Load Breast Cancer dataset
470 breast_cancer = load_breast_cancer()
471 X_bc = breast_cancer.data
472 y_bc = breast_cancer.target.reshape(-1, 1) # Reshape to (n_samples, 1)
473
474 # Split into train and test
475 X_train_bc, X_test_bc, y_train_bc, y_test_bc = train_test_split(
476     X_bc, y_bc, test_size=0.92, random_state=42
477 )
478
479 # Scale features
480 scaler_bc = StandardScaler()
481 X_train_bc = scaler_bc.fit_transform(X_train_bc)
482 X_test_bc = scaler_bc.transform(X_test_bc)
483
484 # Define network architecture for Breast Cancer
485 input_dim_bc = X_train_bc.shape[1]
486 hidden_dim_bc = 16
487 output_dim_bc = 1
488 spectral_dim_hidden_bc = 16 # Adjusted spectral_dim_hidden
489 spectral_dim_output_bc = 1 # Corrected spectral_dim_output to match
output_dim
490
491 layer_dims_bc = [input_dim_bc, hidden_dim_bc, output_dim_bc]
492 spectral_dims_bc = [spectral_dim_hidden_bc, spectral_dim_output_bc]
493 activations_bc = ['relu', 'sigmoid']
494
495 # Initialize the network
496 network_bc = SpectralNeuralNetwork(layer_dims_bc, spectral_dims_bc,
activations_bc)
497
498 # Train the network on Breast Cancer dataset
499 epochs_bc = 500000
500 learning_rate_bc = 0.1 # Adjusted learning rate
501 network_bc.train(X_train_bc, y_train_bc, epochs_bc, learning_rate_bc,
loss_function='bce')

```

```

502 # Plot loss curve for Breast Cancer
503 plot_loss(network_bc.loss_history, "Spectral Neural Network Loss Curve for
504 Breast Cancer Classification")
505
506 # Make predictions on test set
507 predictions_bc = network_bc.predict(X_test_bc)
508 predictions_binary_bc = (predictions_bc > 0.5).astype(int)
509
510 # Calculate MCC
511 mcc_bc = matthews_corrcoef(y_test_bc, predictions_binary_bc)
512
513 # Calculate Accuracy
514 accuracy_bc = accuracy_score(y_test_bc, predictions_binary_bc)
515
516 # Confusion Matrix
517 cm_bc = confusion_matrix(y_test_bc, predictions_binary_bc)
518 plot_confusion_matrix(cm_bc, classes=breast_cancer.target_names, title='
Confusion Matrix for Breast Cancer Classification')
519
520 # ROC Curve
521 # For ROC, we need probabilities or scores
522 # Since our SNN outputs sigmoid activations, we can use them directly
523 roc_scores_bc = predictions_bc.ravel()
524 plot_roc_curve(y_test_bc, roc_scores_bc, title='ROC Curve for Breast Cancer
ROC Curve for Breast Cancer Classification')
525
526 # =====
527 # Part 3: Regression on Diabetes Dataset
528 # =====
529 print("\n\nTraining on Diabetes Regression")
530
531 # Load Diabetes dataset
532 diabetes = load_diabetes()
533 X_diab = diabetes.data
534 y_diab = diabetes.target.reshape(-1, 1) # Reshape to (n_samples, 1)
535
536 # Split into train and test
537 X_train_diab, X_test_diab, y_train_diab, y_test_diab = train_test_split(
538     X_diab, y_diab, test_size=0.92, random_state=42
539 )
540
541 # Scale features
542 scaler_diab = StandardScaler()
543 X_train_diab = scaler_diab.fit_transform(X_train_diab)
544 X_test_diab = scaler_diab.transform(X_test_diab)
545
546 # Define network architecture for Diabetes
547 input_dim_diab = X_train_diab.shape[1]
548 hidden_dim_diab = 16
549 output_dim_diab = 1
550 spectral_dim_hidden_diab = 16 # Adjusted spectral_dim_hidden
551 spectral_dim_output_diab = 1 # Corrected spectral_dim_output to match
output_dim
552
553 layer_dims_diab = [input_dim_diab, hidden_dim_diab, output_dim_diab]
554 spectral_dims_diab = [spectral_dim_hidden_diab, spectral_dim_output_diab]
555 activations_diab = ['relu', 'relu'] # Using 'relu' for regression output
556
557 # Initialize the network
558 network_diab = SpectralNeuralNetwork(layer_dims_diab, spectral_dims_diab,
activations_diab)
559

```

```

560 # Train the network on Diabetes dataset
561 epochs_diab = 1000000
562 learning_rate_diab = 0.005 # Adjusted learning rate for regression
563 network_diab.train(X_train_diab, y_train_diab, epochs_diab,
                    learning_rate_diab, loss_function='mse')
564
565 # Plot loss curve for Diabetes
566 plot_loss(network_diab.loss_history, "Spectral Neural Network Loss Curve
for Diabetes Regression")
567
568 # Make predictions on test set
569 predictions_diab = network_diab.predict(X_test_diab)
570
571 # Calculate R2 Score
572 r2_diab = r2_score(y_test_diab, predictions_diab)
573
574 # Scatter Plot for Regression
575 plot_regression_predictions(y_test_diab, predictions_diab, "Predicted vs
Actual Values for Diabetes Regression")
576
577 # =====
578 # Part 4: Comparing with Traditional ML Models
579 # =====
580
581 # =====
582 # Classification Comparison
583 # =====
584 print("\n\nComparing Classification Models on Breast Cancer Dataset")
585
586 # Initialize other classification models
587 svm_bc = SVC(probability=True, random_state=42)
588 mlp_bc = MLPClassifier(hidden_layer_sizes=(16,), max_iter=1000,
random_state=42)
589 logistic_bc = LogisticRegression(max_iter=1000, random_state=42)
590 rf_bc = RandomForestClassifier(n_estimators=100, random_state=42)
591
592 # Train models
593 svm_bc.fit(X_train_bc, y_train_bc.ravel())
594 mlp_bc.fit(X_train_bc, y_train_bc.ravel())
595 logistic_bc.fit(X_train_bc, y_train_bc.ravel())
596 rf_bc.fit(X_train_bc, y_train_bc.ravel())
597
598 # Make predictions
599 predictions_svm_bc = svm_bc.predict(X_test_bc)
600 predictions_mlp_bc = mlp_bc.predict(X_test_bc)
601 predictions_logistic_bc = logistic_bc.predict(X_test_bc)
602 predictions_rf_bc = rf_bc.predict(X_test_bc)
603
604 # Compute MCC for all models
605 mcc_svm_bc = matthews_corrcoef(y_test_bc, predictions_svm_bc)
606 mcc_mlp_bc = matthews_corrcoef(y_test_bc, predictions_mlp_bc)
607 mcc_logistic_bc = matthews_corrcoef(y_test_bc, predictions_logistic_bc)
608 mcc_rf_bc = matthews_corrcoef(y_test_bc, predictions_rf_bc)
609
610 # Spectral Neural Network MCC already computed as mcc_bc
611
612 # Prepare data for comparison
613 models_classification = ['Spectral NN', 'SVM', 'MLP', 'Logistic Regression',
, 'Random Forest']
614 mcc_scores_classification = [mcc_bc, mcc_svm_bc, mcc_mlp_bc,
mcc_logistic_bc, mcc_rf_bc]
615
616 # Plot MCC comparison

```

```

617 plot_classification_comparison(
618     mcc_scores_classification,
619     models_classification,
620     'MCC Comparison on Breast Cancer Classification'
621 )
622
623 # =====
624 # Regression Comparison
625 # =====
626 print("\n\nComparing Regression Models on Diabetes Dataset")
627
628 # Initialize other regression models
629 lr_diab = LinearRegression()
630 rf_diab = RandomForestRegressor(n_estimators=100, random_state=42)
631
632 # Train models
633 lr_diab.fit(X_train_diab, y_train_diab.ravel())
634 rf_diab.fit(X_train_diab, y_train_diab.ravel())
635
636 # Make predictions
637 predictions_lr_diab = lr_diab.predict(X_test_diab).reshape(-1,1)
638 predictions_rf_diab = rf_diab.predict(X_test_diab).reshape(-1,1)
639
640 # Compute R2 Score for all models
641 r2_lr_diab = r2_score(y_test_diab, predictions_lr_diab)
642 r2_rf_diab = r2_score(y_test_diab, predictions_rf_diab)
643 # Spectral Neural Network R2 already computed as r2_diab
644
645 # Prepare data for comparison
646 models_regression = ['Spectral NN', 'Linear Regression', 'Random Forest']
647 r2_scores_regression = [r2_diab, r2_lr_diab, r2_rf_diab]
648
649 # Plot R2 comparison
650 plot_regression_comparison(
651     r2_scores_regression,
652     models_regression,
653     'R2 Score Comparison on Diabetes Regression'
654 )
655
656 # Scatter Plot Comparisons
657 plt.figure(figsize=(8,6))
658 plt.scatter(y_test_diab, predictions_diab, alpha=0.5, label='Spectral NN')
659 plt.scatter(y_test_diab, predictions_lr_diab, alpha=0.5, label='Linear
Regression')
660 plt.scatter(y_test_diab, predictions_rf_diab, alpha=0.5, label='Random
Forest')
661 plt.plot([y_test_diab.min(), y_test_diab.max()], [y_test_diab.min(),
y_test_diab.max()], 'k--', lw=2, label='Ideal Fit')
662 plt.xlabel('Actual Values')
663 plt.ylabel('Predicted Values')
664 plt.title('Predicted vs Actual Values for Diabetes Regression')
665 plt.legend()
666 plt.grid(True)
667 plt.show()
668
669 # =====
670 # Summary of Results
671 # =====
672 print("\n\nSummary of Classification Results on Breast Cancer Dataset:")
673 for model, mcc in zip(models_classification, mcc_scores_classification):
674     print(f"{model}: MCC = {mcc:.4f}")
675
676 print("\n\nSummary of Regression Results on Diabetes Dataset:")

```



```
677     for model, r2 in zip(models_regression, r2_scores_regression):
678         print(f"{model}: R2 Score = {r2:.4f}")
679
680 if __name__ == "__main__":
681     main()
```

Listing 1: Spectral Neural Networks Implementation